


BY:

NATHAN DEKEYREL  
AND WADE FORTSON

# ANALYSIS OF THE COMPLETELY FAIR SCHEDULER



# GOALS OF PROJECT



We wanted to provide an analysis of an actual scheduling algorithm used in modern operating systems



While researching scheduling algorithms, we found the Completely Fair Scheduler used in Linux



We wanted implement CFS to compare our implementation and Linux's

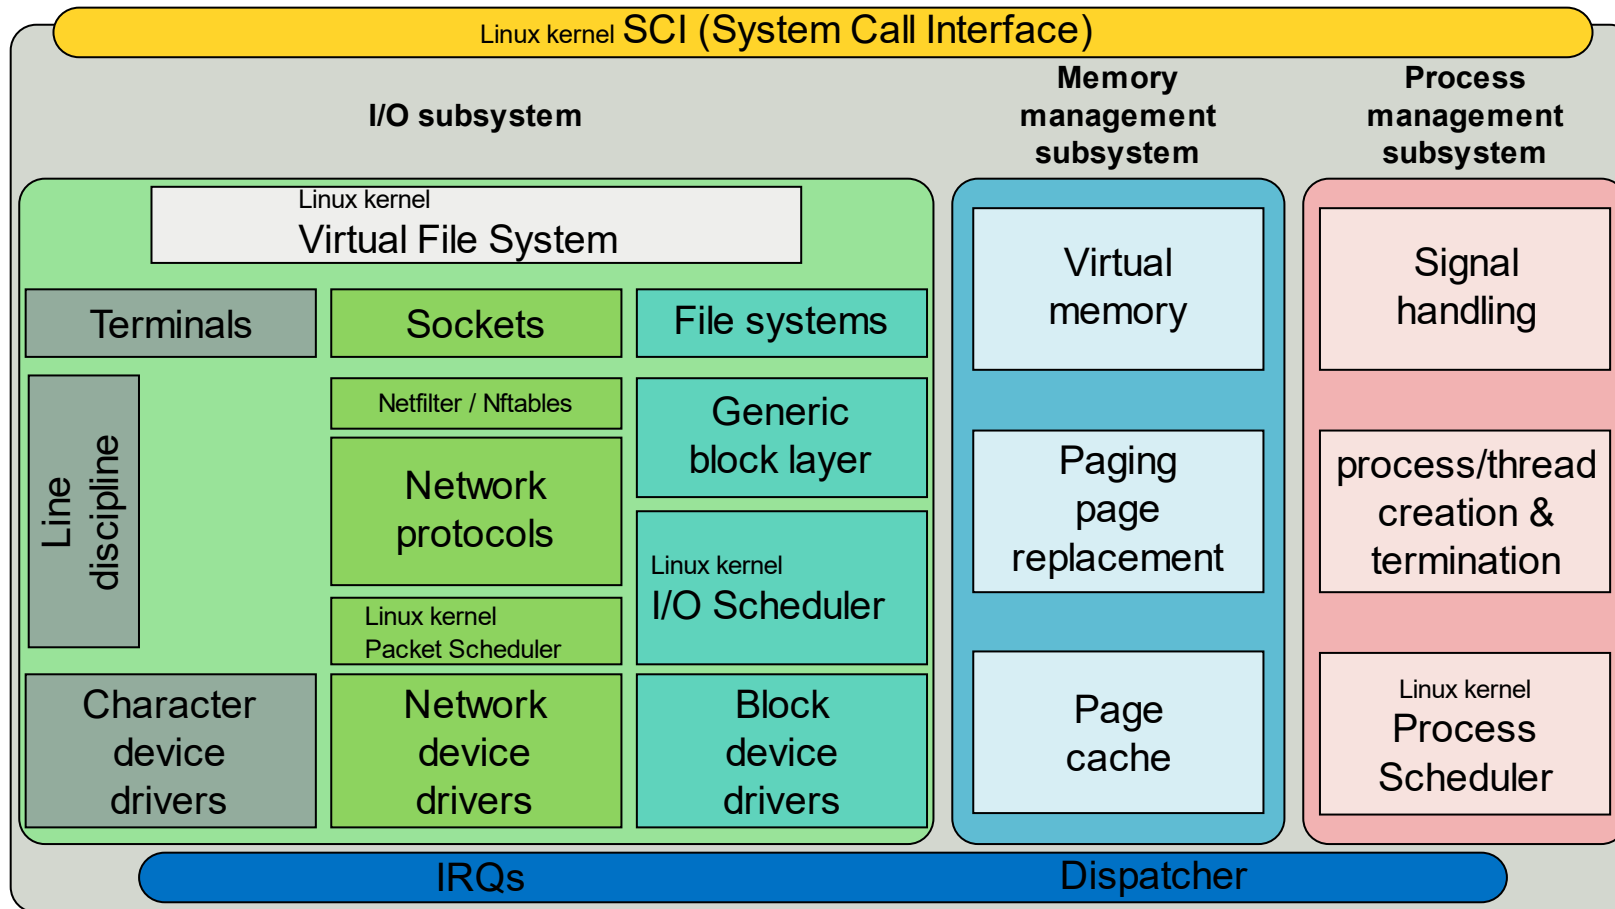
## BACKGROUND

In class we saw 6 different scheduling algorithms

- FCFS, RR, SPN, SRT, HRRN, and FB

With a good understanding of these, we wanted to see what's used in industry

- Seeing how our own implementation compares



# SIMPLIFIED LINUX DIAGRAM

# CFS IN LINUX

The main scheduler used by the Linux kernel since 2007



```
graph TD; A[The main scheduler used by the Linux kernel since 2007] --> B[Uses Red-black tree for sorting – O(log n) insertion and deletion]; B --> C[Inspired by 'Ideal Fair Queueing' in network packet scheduling]; C --> D[Replaced the O(1) scheduler previously used];
```

Uses Red-black tree for sorting –  $O(\log n)$  insertion and deletion

Inspired by 'Ideal Fair Queueing' in network packet scheduling


Replaced the  $O(1)$  scheduler previously used

# O(1) SCHEDULER



## Did not contain any algorithms that ran worse than $O(1)$ time

Every part of scheduler guaranteed to complete in an upper-bounded, definite time regardless of the size of input



## Algorithm

Uses 2 arrays – Active and Expired

Each process given fixed time quantum

- After exhausted, preempted and moved to Expired

Once active emptied, swap arrays (using pointers) and repeat

# ISSUES WITH $O(1)$ SCHEDULER



## Identifies interactivity based on average sleep time

Processes that wait for user input could be assumed inactive

- Gives priority to interactive tasks
- Penalizing non-interactive by lowering priority

All calculations to determine interactivity are prone to miscalculation



## High complexity

Extremely complicated algorithms with many edge cases




## Unpredictable

Prioritizing response to user input can cause random shift in ordering

# DESIGN OF CFS

- Goals
  - Reduce complexity – No special cases
  - Fair CPU distribution across processes
  - Better responsiveness – no misclassification of tasks
  - Predictable





# CFS GROUNDWORK

Task – minimal entity that can be scheduled


- Can also be a group of threads, leading to schedulable entities

Sorts schedulable entities in tree based on time

- Minimum vruntime stored in leftmost node
- Maximum execution time based on ‘ideal processor’
  - Calculated as  $\text{time\_waiting} / \text{total\_num\_of\_processes}$

# RB TREE RECAP

- Every node is either Red or Black
  - Root is black
  - Leaves (which are NIL) are black
  - Red nodes can't have red children
  - All paths have the same number of black nodes
- Guarantees  $O(\log n)$  for Insert, Delete, and finding minimum node
  - Also, self balancing



# RB TREE FUNCTIONS

## Helper functions:

- Rotate left
- Rotate right
- Handle Node Removal Issues

## Tree Operations

- Insert
- Delete

## Tree Traversal

- First
- Last
- Next
- Prev

# IMPLEMENTATION DECISIONS

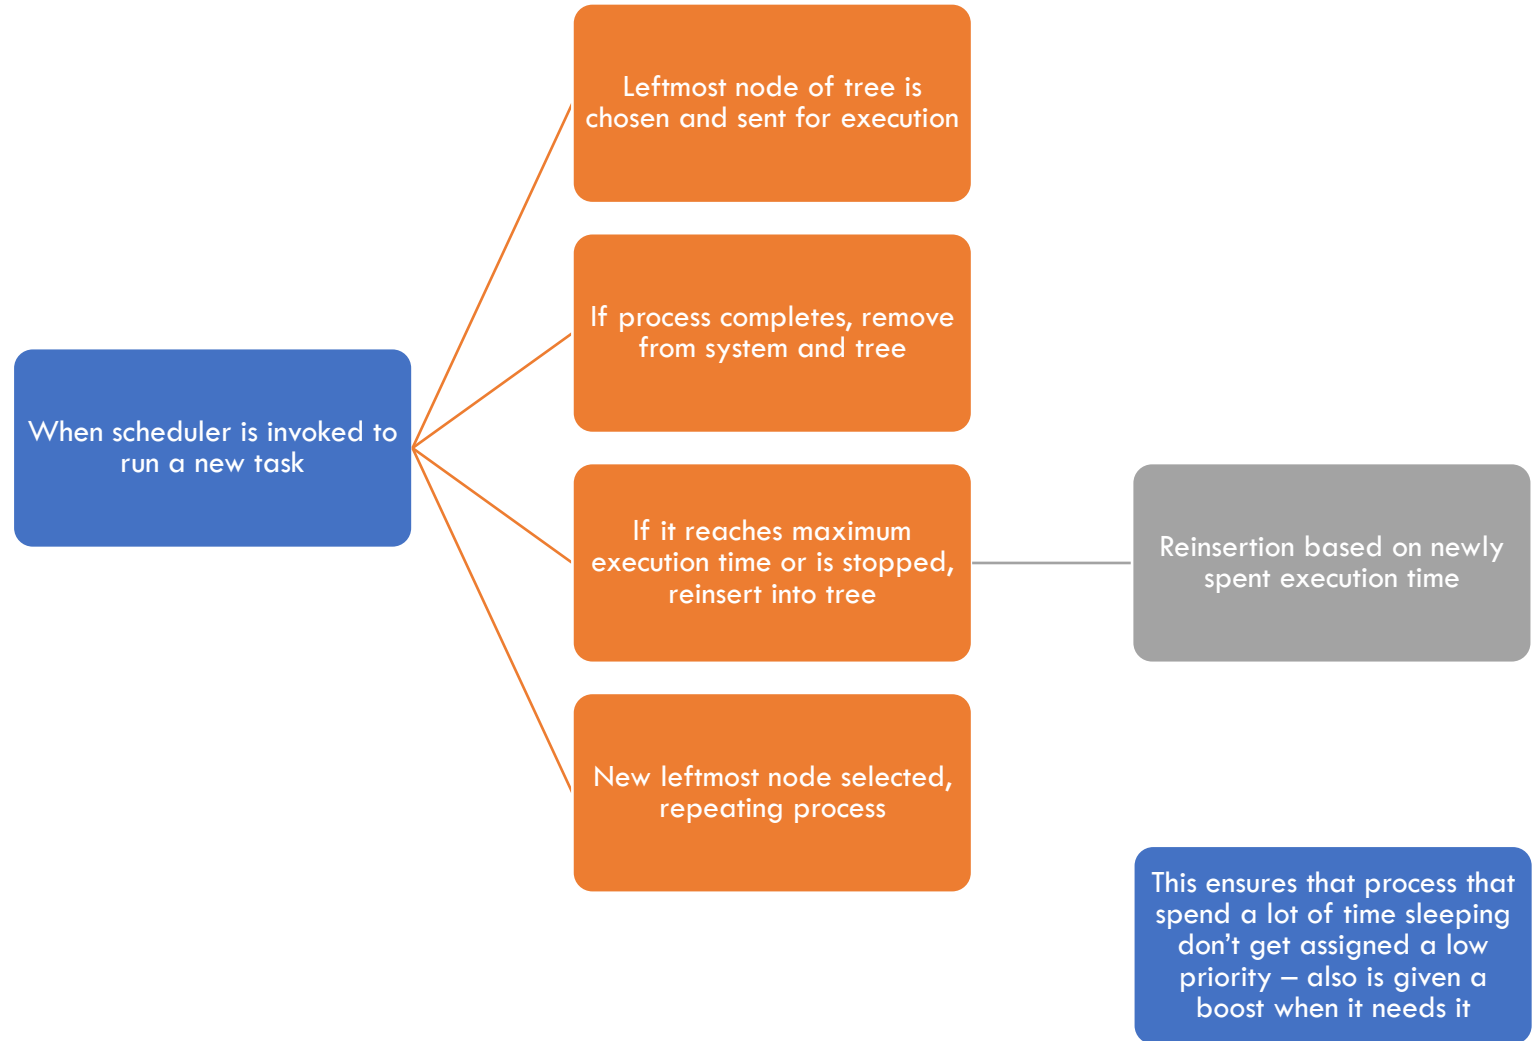
## Linux

- Handles concurrency
- Advanced memory management
- Multiple tree traversal strategies
- Production level code

## Us

- Not concurrent
- Simple memory management
- Only uses one traversal strategy
- Only suitable for simulation

# CFS ALGORITHM



# CFS V O(1)

## CFS

- CFS
  - Pros
    - Predictable
    - Simple
    - Good for interactivity
  - Cons
    - Higher time complexity
    - Can lead to latency spikes
    - Expensive if lots of tasks
    - Memory overhead due to vruntime

## O(1)

- O(1)
  - Pros
    - Defined upper-bound runtime regardless of workload size
  - Cons
    - Unpredictable
    - Subject to miscalculations
    - Low responsiveness

“CFS Basically models an ‘ideal, precise multitasking CPU’ on real hardware” – Ingo Molnar

# PROGRAM DESIGN



# SYSTEM ARCHITECTURE

1

## Initialize workloads

1. Using python, we initialize workloads with semi-random values
2. These are generated as CSVs

2

## Parse csv

- Parse the CSVs into a queue of process structures
- Containing the necessary info for our scheduler

3

## Design the CFS algorithm

1. Using a red-black tree for scheduling operations



# PROCESS DESIGN

We created a general process struct holding information relevant to each scheduler

We categorized attributes into

Basic

Timing

CFS required

State

Queue Management

# PROCESS DESIGN

## Base attributes needed

Process id

priority



Timing – contains necessary attributes for scheduling and simulation

Arrival time

First run time

Completion time

Burst time

Total time

# PROCESS DESIGN



CFS specific attributes

vruntime  
rb\_node



State

We decided to use an enum for

- NEW, READY, RUNNING, WAITING, and COMPLETED processes



Queue Management

Keep track of next and previous process in queue

# METRICS



Response Time

first runtime – arrival time



Turnaround Time

completion time – start time



Virtual Runtime Variance

$\text{delta\_exec\_weighed} = \text{delta\_exec} * (\text{NICE\_0\_LOAD} / \text{curr-}\>\text{load.weight})$

**SHOW CODE**

Generally, the drawbacks of CFS aren't noticeable

- For most cases, the default CFS algorithm works well enough
  - In modern OS, there usually aren't enough tasks to cause performance decreases
  - For specialized cases, tuning options are available within the Linux kernel
- For real-time tasks, the Linux kernel makes 2 separate schedulers available
  - SCHED\_RR and SCHED\_FIFO

If you have a highly specialized system, you can make a scheduler to fit your needs as the Linux kernel is modular

# FINAL THOUGHTS